FGV EMAp V-DS FAPESP SÃO PAULO RESEARCH FOUNDATION

Assessing Urban Environments with Vision-Language Models

A **Comparative Analysis** of **AI-Generated Ratings** and **Human Evaluations**

Felipe Moreno-Vera and Jorge Poco

Fundação Getulio Vargas / Visual Data Science Lab

www.visualdslab.com

Context & Motivation

Techniques to analyze the urban perception

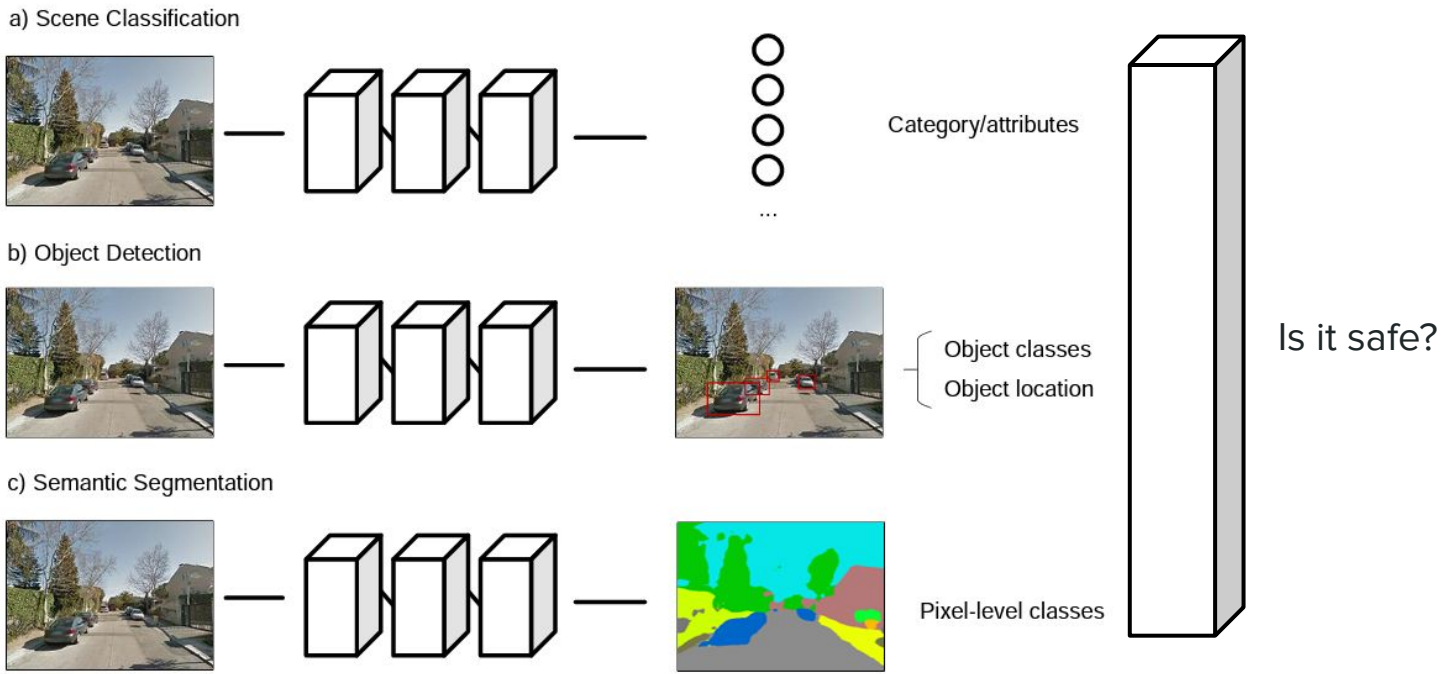

Hint: Street view images **contain rich, complex scenes** that are **difficult** to **interpret** purely through **pixels** or vector **embeddings**.

What information is obtained?

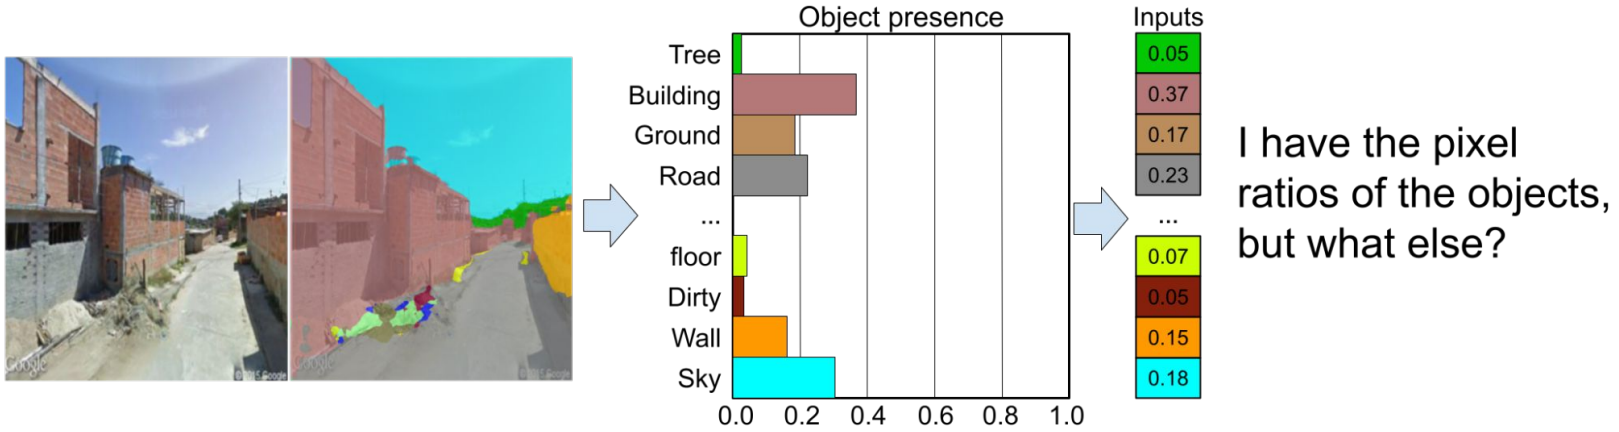

- Traditional methods **identify objects** – e.g., they see "a building, sidewalks, roads"

- But they may **miss context** — e.g., they see “a **building**” but not "a **neglected building** with **graffiti** and **broken windows**."

How to obtain the rich information (context) from images?

The street's edges are **poorly defined**. On the right, there is a narrow, **rudimentary** concrete **sidewalk** and **minor cracks** visible in the **unfinished wall**. On the left, there is **no sidewalk** at all; the pavement gives way to dirt, sand, and piles of construction debris that spill onto the road. Besides, It is **littered** with **construction materials**, sand, and other **rubble**.

We have a description about the current status of the street

- Image descriptions adds **semantic depth** and **human-like perception**.
- Vision-Language models analyze the **context of the image** and provide further information.

Overview

- Place Pulse
- Image-to-Text descriptions
- UrbanVLM
- Ablation study
- Conclusions

Place Pulse

Place Pulse dataset

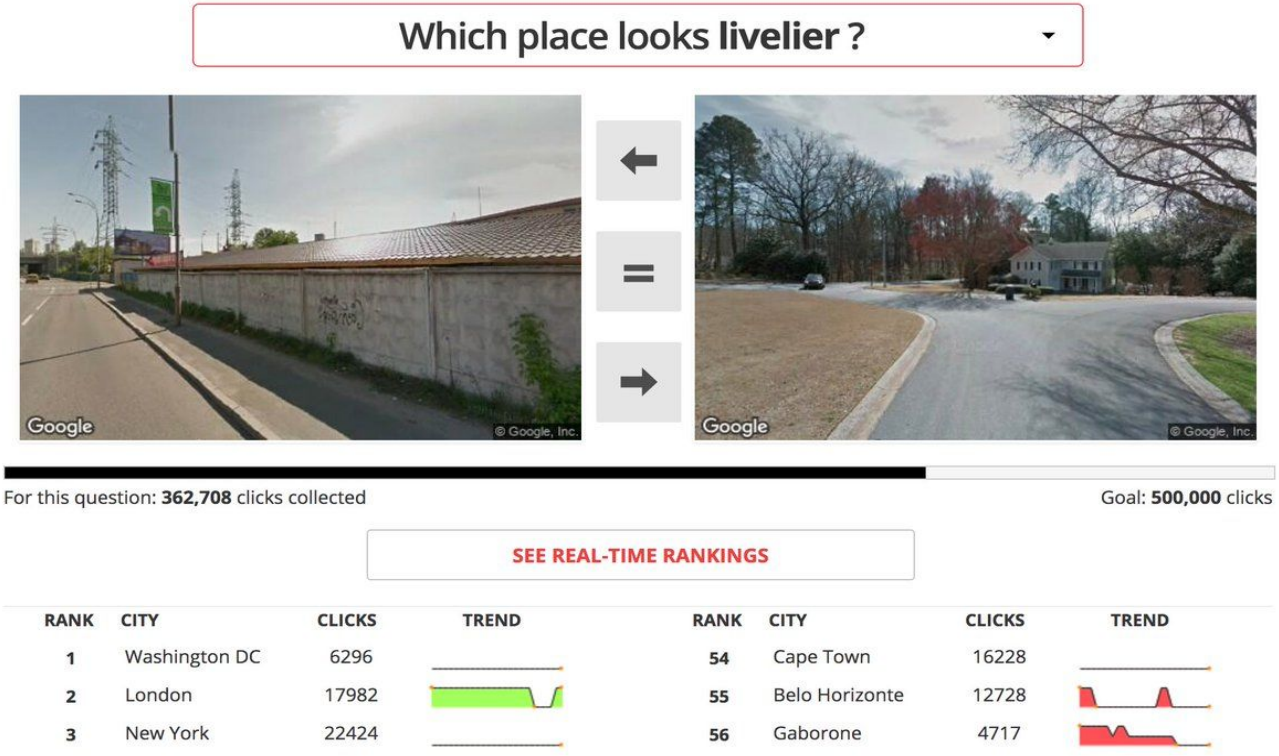

https://centerforcollectivelearning.org/urbanperception

Cities included

- 1 223 649 Comparisons
- 111 390 images
- 32 countries
- 56 cities
- 6 categories:
  - Safety
  - Boring
  - Depressing
  - Wealthy
  - Lively
  - Beauty

Note: Same color means same country.

Strength of Schedule

left	right	winner
		draw
		left
		right
. . .	. . .	. . .
		right
		left

Image | Perceptual Scores

High safety scores images

Low safety scores images

Image-to-Text descriptions

Types of descriptions

- **General description:** Simple description of the image
- **Subjective description:** Based on the perception
  - **Positive:** Describe the image based on its corresponding perception.
  - **Negative:** Describe the image based on its opposite perception.

Image general descriptions

Image		
Model	Description	Description
LlaVA	The image depicts a narrow alleyway between two buildings, with one of the buildings being a brick structure. The alleyway is surrounded by a dirt road, and there are a few cars parked along the road.	The image shows a residential area with a well-maintained hedge around a house and several potted plants, creating a pleasant, aesthetic, and inviting atmosphere.
BLIP-2	This image shows a narrow street in a residential area under development or construction. The buildings are primarily made of exposed and unfinished red bricks and concrete.	The image shows a residential street scene. Additionally, a tall hedge covers a gate and wall, possibly concealing a private residence.
BLIP	This is a Google Street View image of a building under construction.	This is a Google Street View image of a green residential area in the Philippines.

Image general descriptions

Image		
Model	Description	Description
LlaVA	The image depicts a narrow alleyway between two buildings, with one of the buildings being a brick structure. The alleyway is surrounded by a dirt road, and there are a few cars parked along the road.	The image shows a residential area with a well-maintained hedge around a house and several potted plants, creating a pleasant, aesthetic, and inviting atmosphere.
BLIP-2	This image shows a narrow street in a residential area under development or construction. The buildings are primarily made of exposed and unfinished red bricks and concrete.	The image shows a residential street scene. Additionally, a tall hedge covers a gate and wall, possibly concealing a private residence.
[illegible]	[illegible] Google Street View image of [illegible] construction.	This is a Google [illegible] green residential [illegible]

Randomly select 50 samples and compare the description results

Urban Vision-Language Model

Model architecture

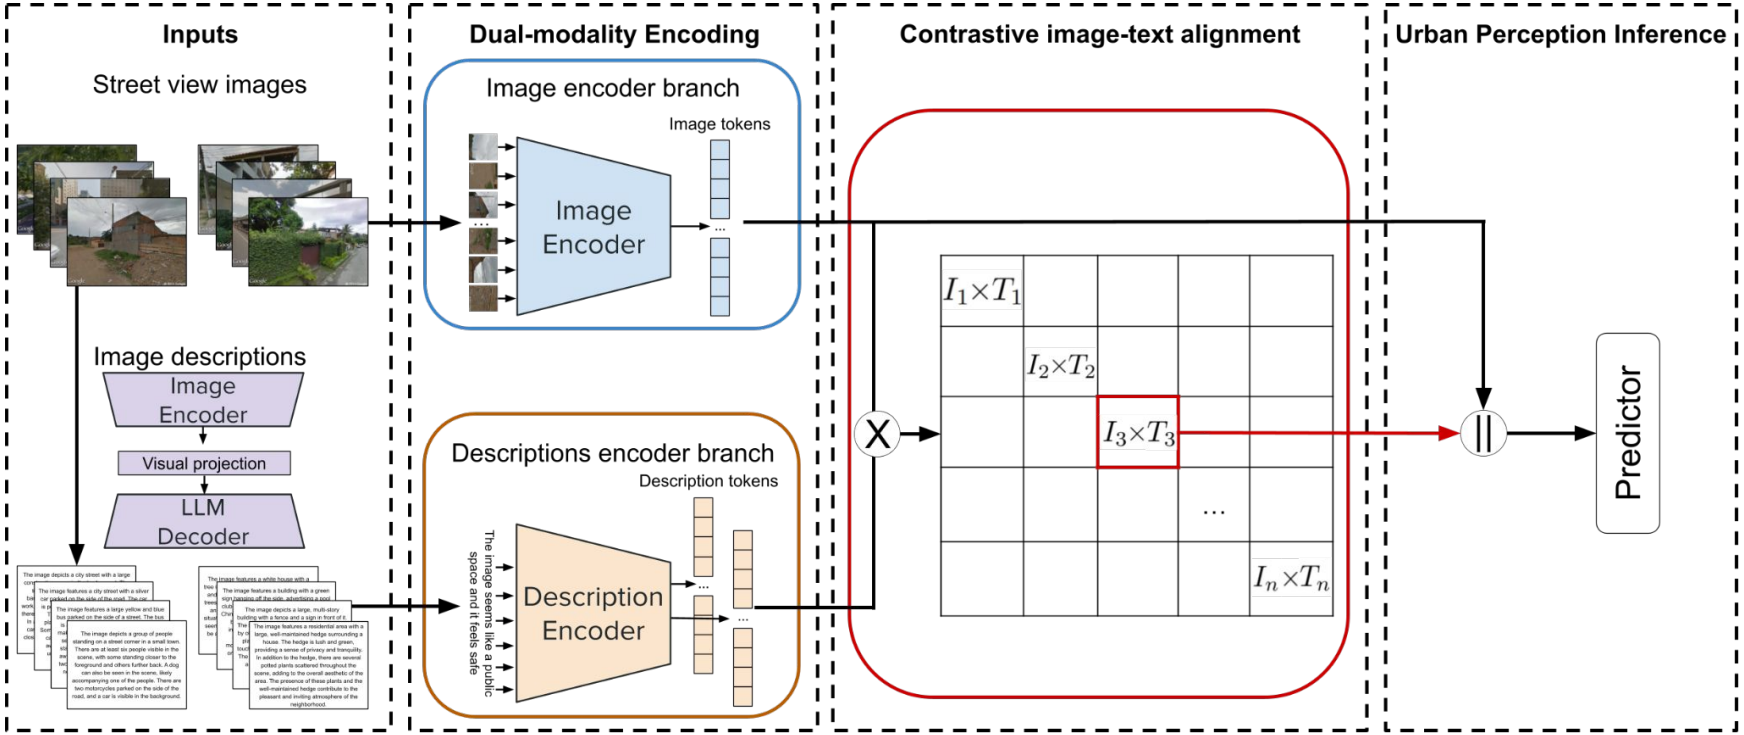

Model architecture - image descriptions

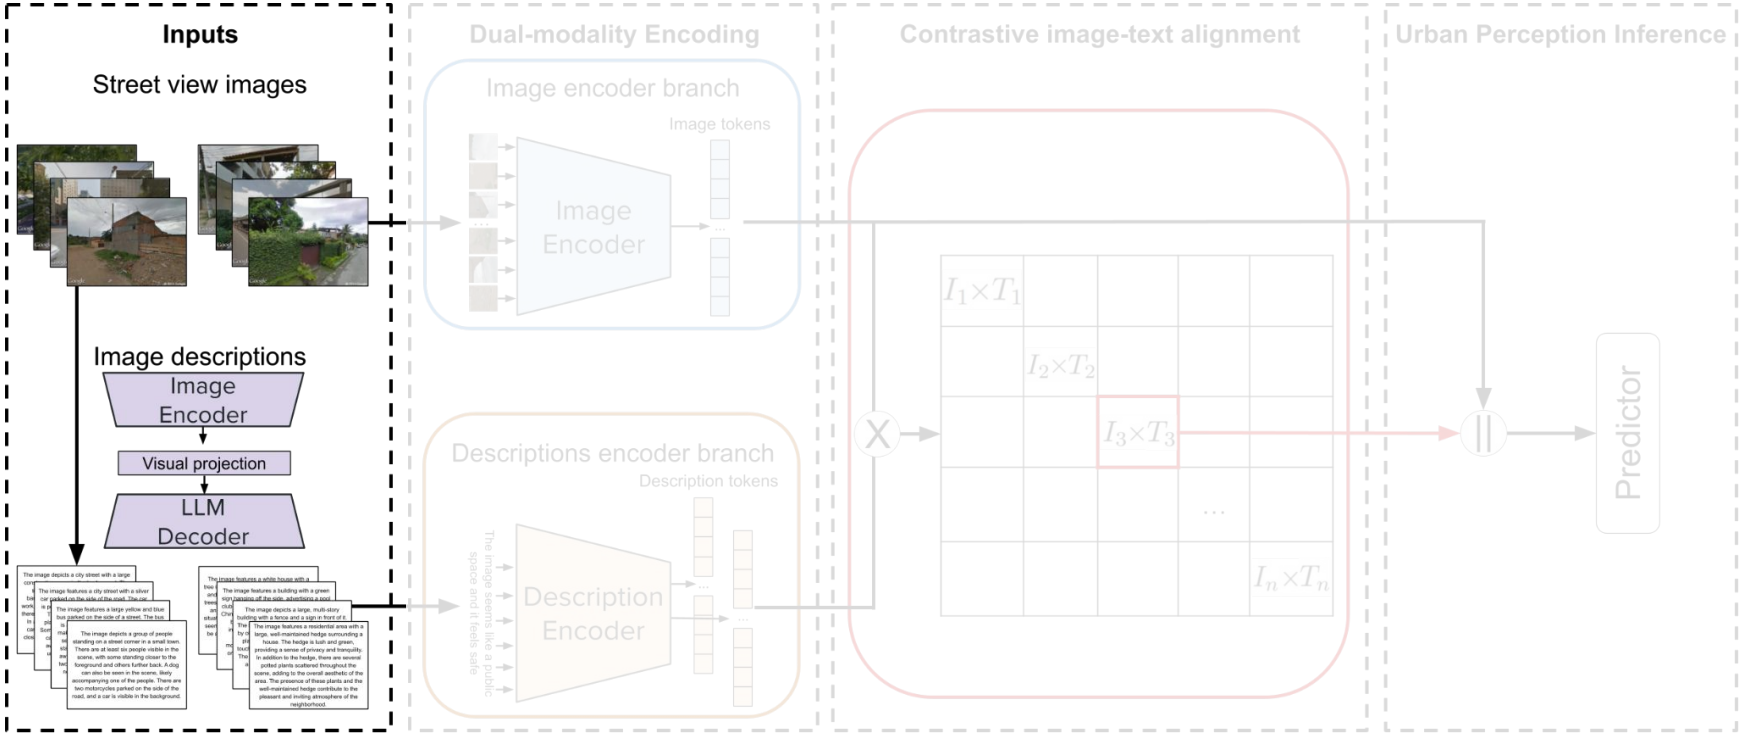

We use and compares **LlaVA** and **BLIP-2** performances

Model architecture - image-text encoders

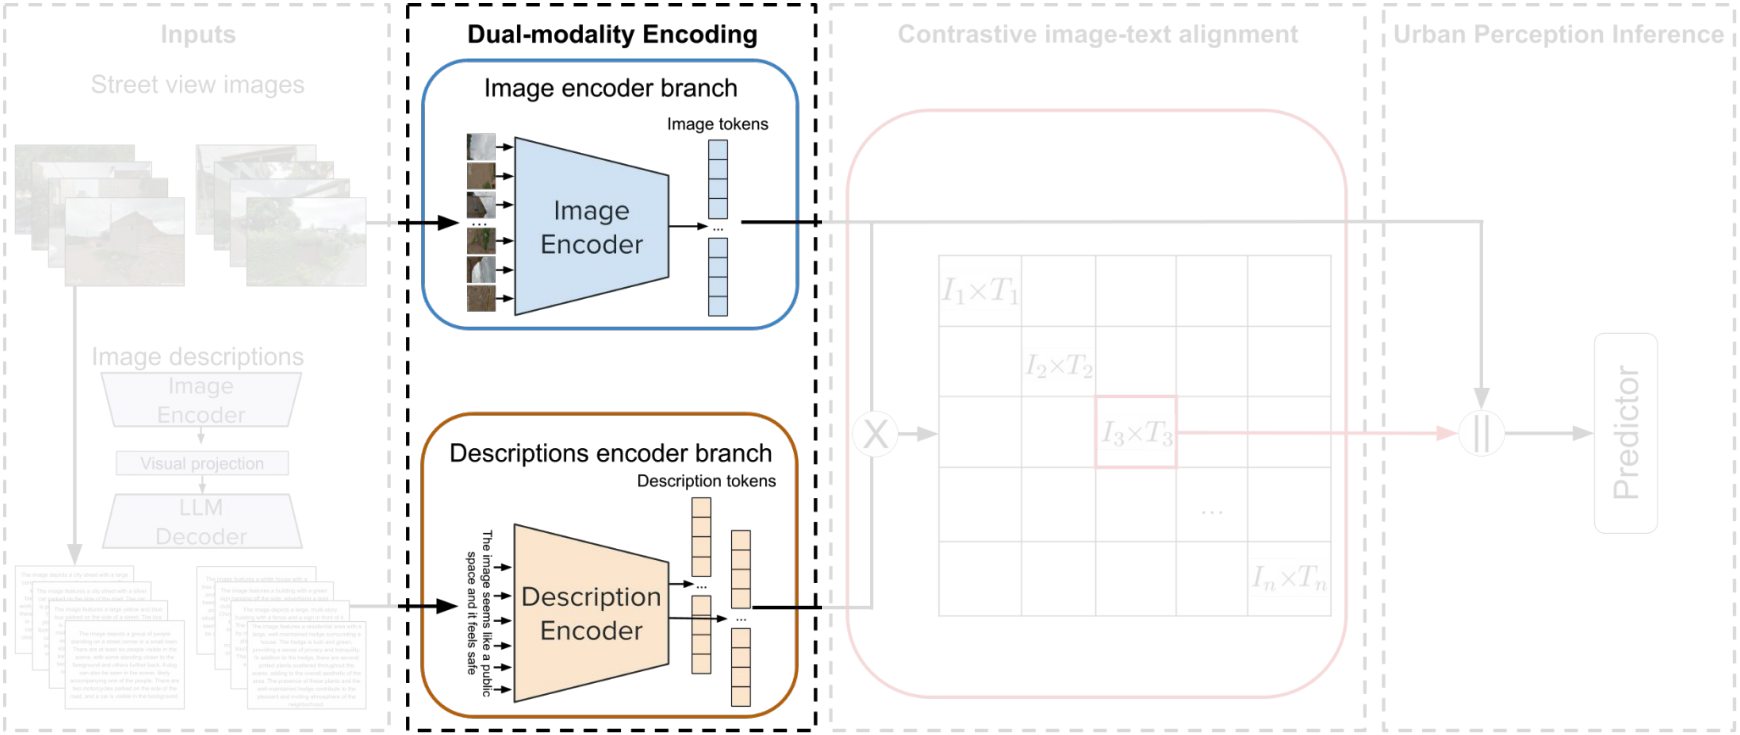

We use and compares **CLIP** and **SigLIP** performances

Model architecture - contrastive

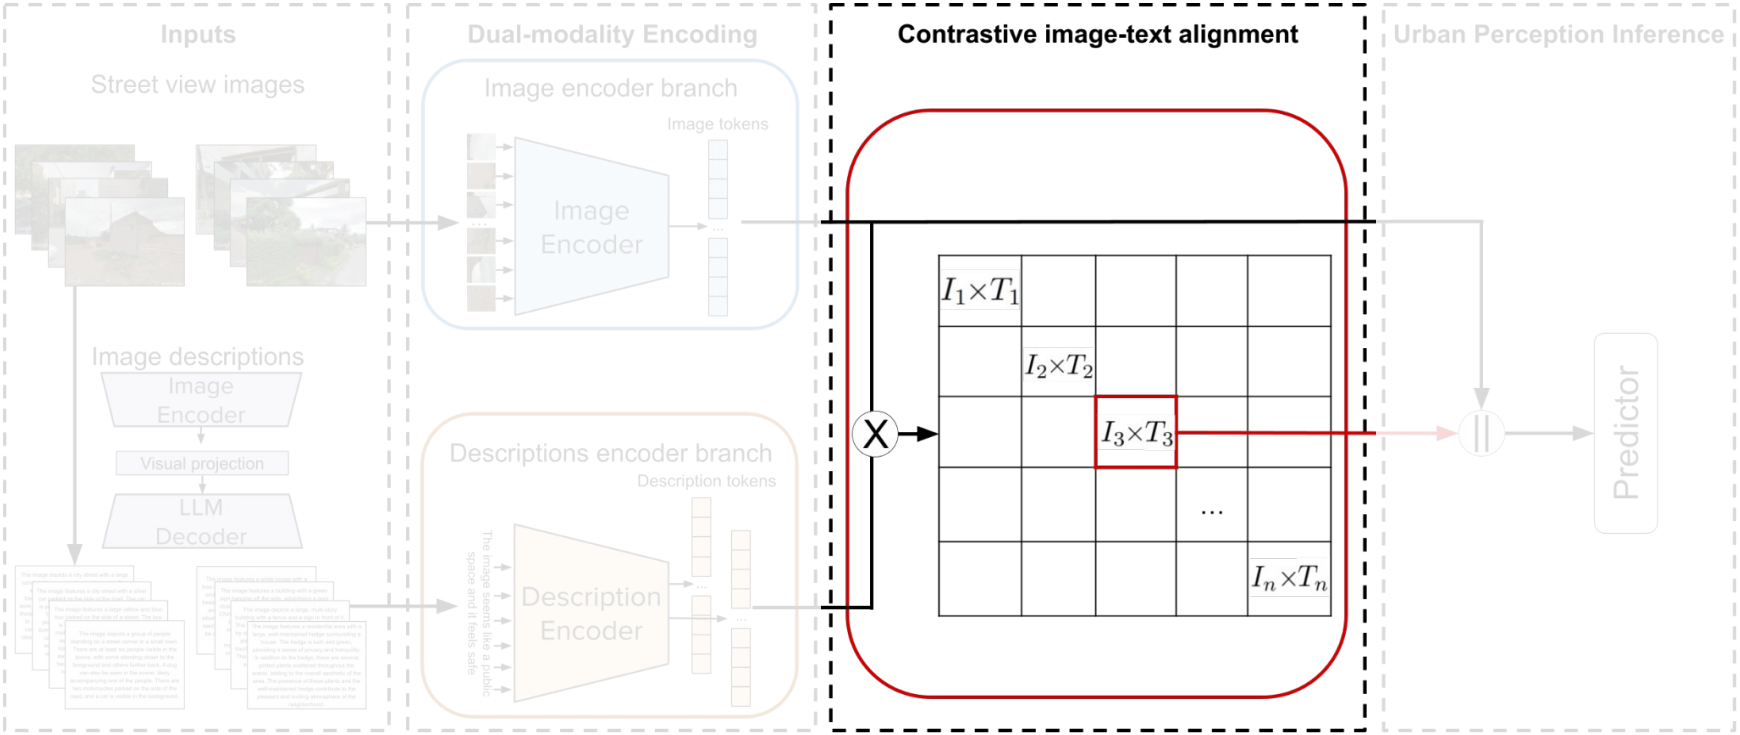

Model architecture - heads

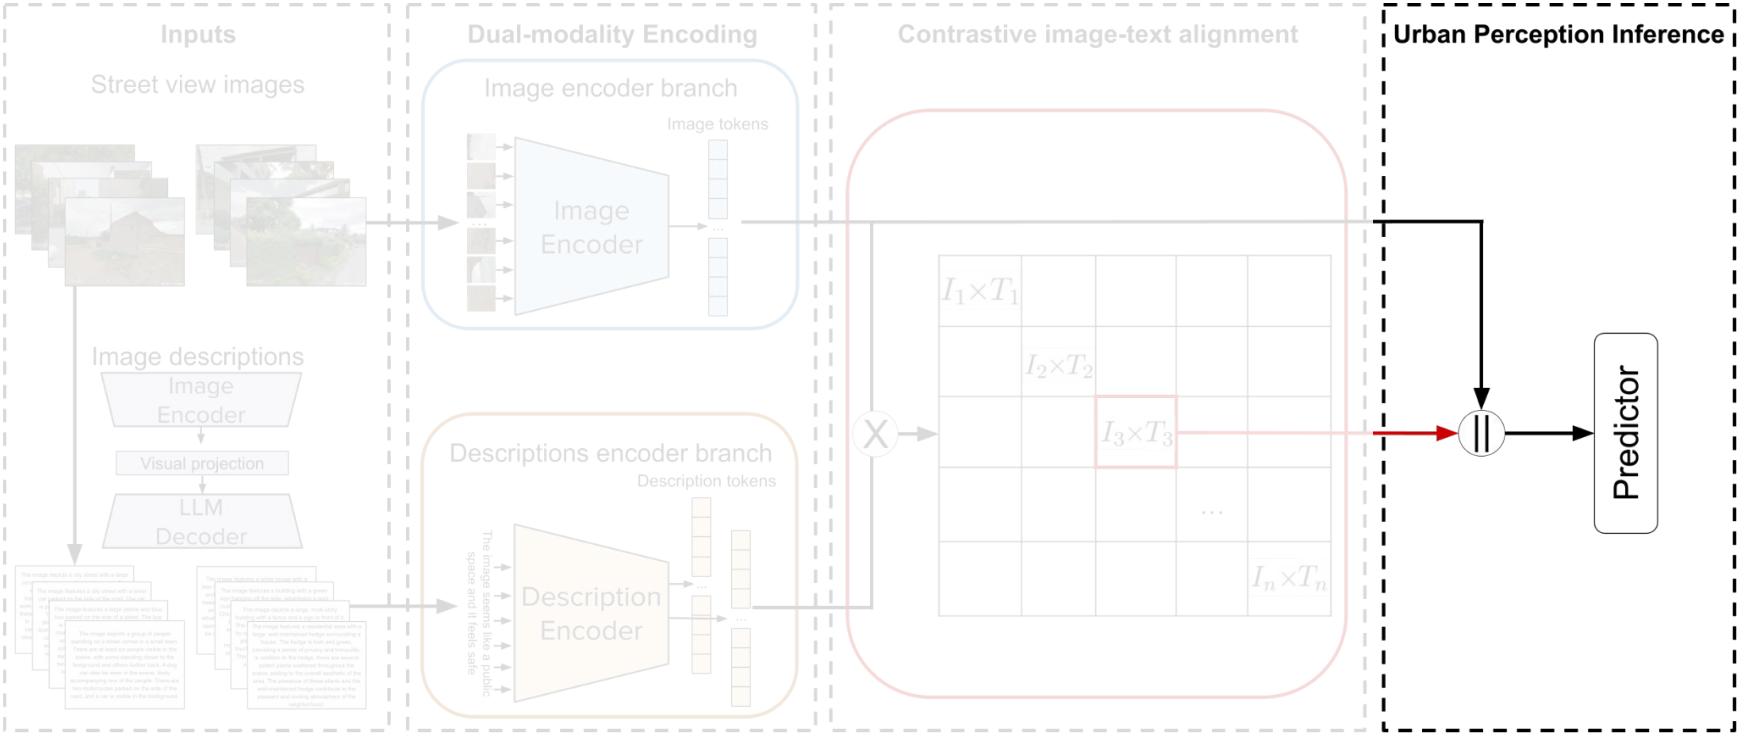

Classification and regression results

Classification

Model	Acc
PspNet+VGG [29]	48.38
DeepLabV3+VGG [29]	51.93
DSAPN+ResNet [54]	64.87
MTDRALN-LC [25]	65.07
MTDRALN-TC [25]	65.82
VGG+ImageNet [28]	65.72
VGG-GAP+ImageNet [28]	66.09
VGG+Places365 [28]	66.46
VGG-GAP+Places365 [28]	66.96
VGG19+ImageNet [4]	67.01
PSPNet+SVR [55]	70.63
DeiT+ResNet50 [40]	71.16
ViT-nn [27]	71.29
ViT-nn+OneFormer [27]	75.68
UrbanVLM (LlaVA+SigLIP)	82.55

Regression

Model	R^2	RMSE
PSPNet-Regressor [55]	0.25	–
Fine-Tuned BERT [22]	0.42	–
FPN-based regressor [20]	0.52	–
DeepLabV3+ regressor [20]	–	2.16
DeepLabV3+ regressor [52]	–	2.91
SFB5+ConvNeXt-B+RF [60]	0.67	1.29
VIT+SegFormer+RF [11]	0.76	1.75
UrbanVLM (LlaVA+CLIP)	0.84	1.04

Ablation study

Components

We define 4 main components:

- **Heads:** Classification and regression MLP
- **Dual-modality:** Linear projections from Image and Text encoders
- **Image-to-Text:** Generates **positive** and **negative** descriptions
- **Contrastive Learning**: Image-text alignment

Only heads *(learns to classify and regress)*

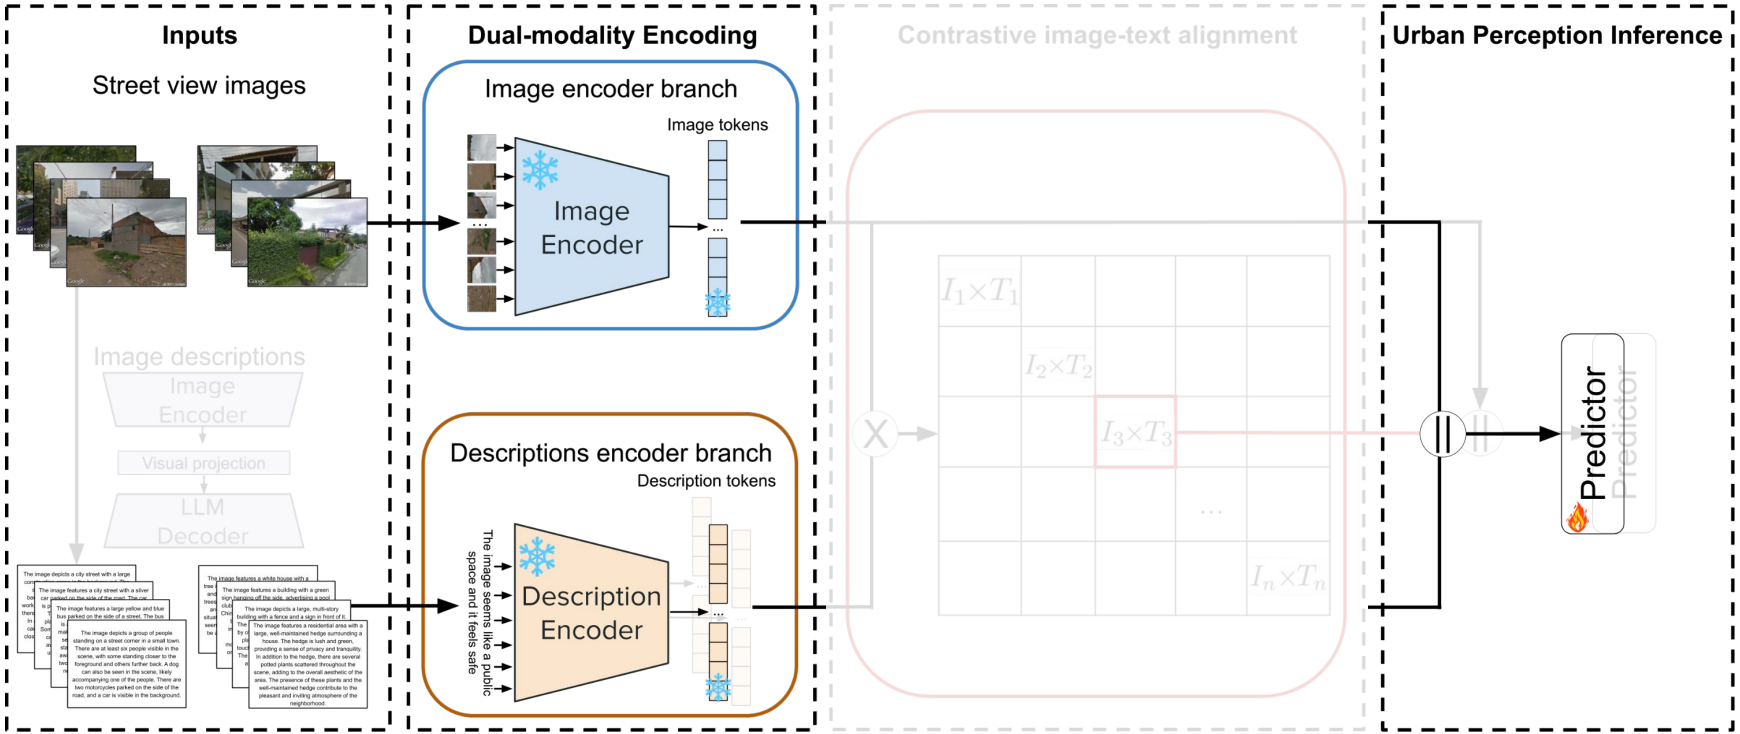

Use the corresponding **positive** description and concatenate.

Dual-modality *(learns to project image and text)*

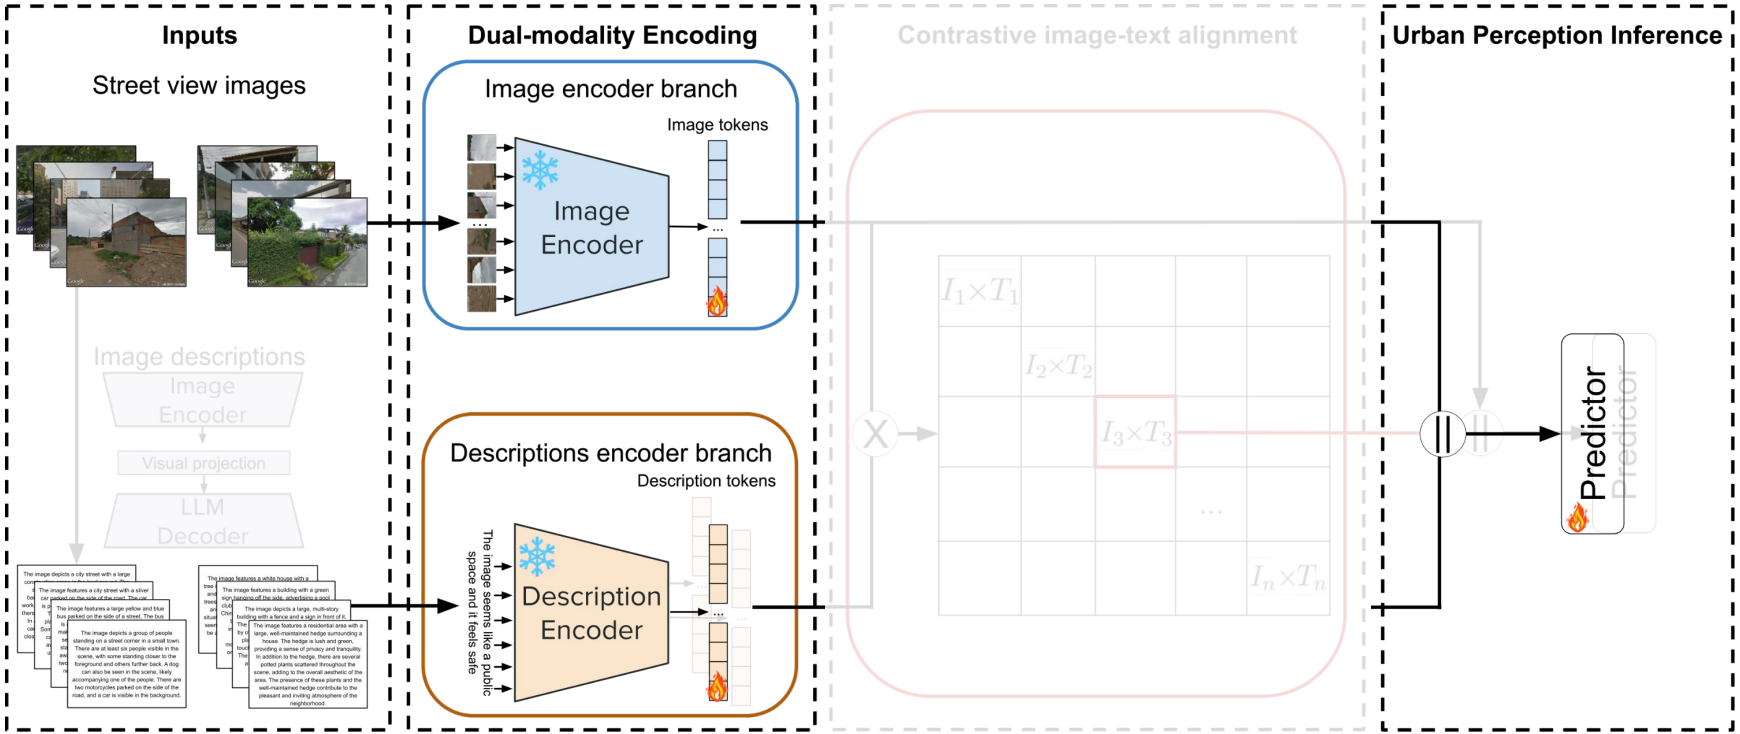

Use the corresponding **positive** description and concatenate.

Image-to-text *(learns to describe)*

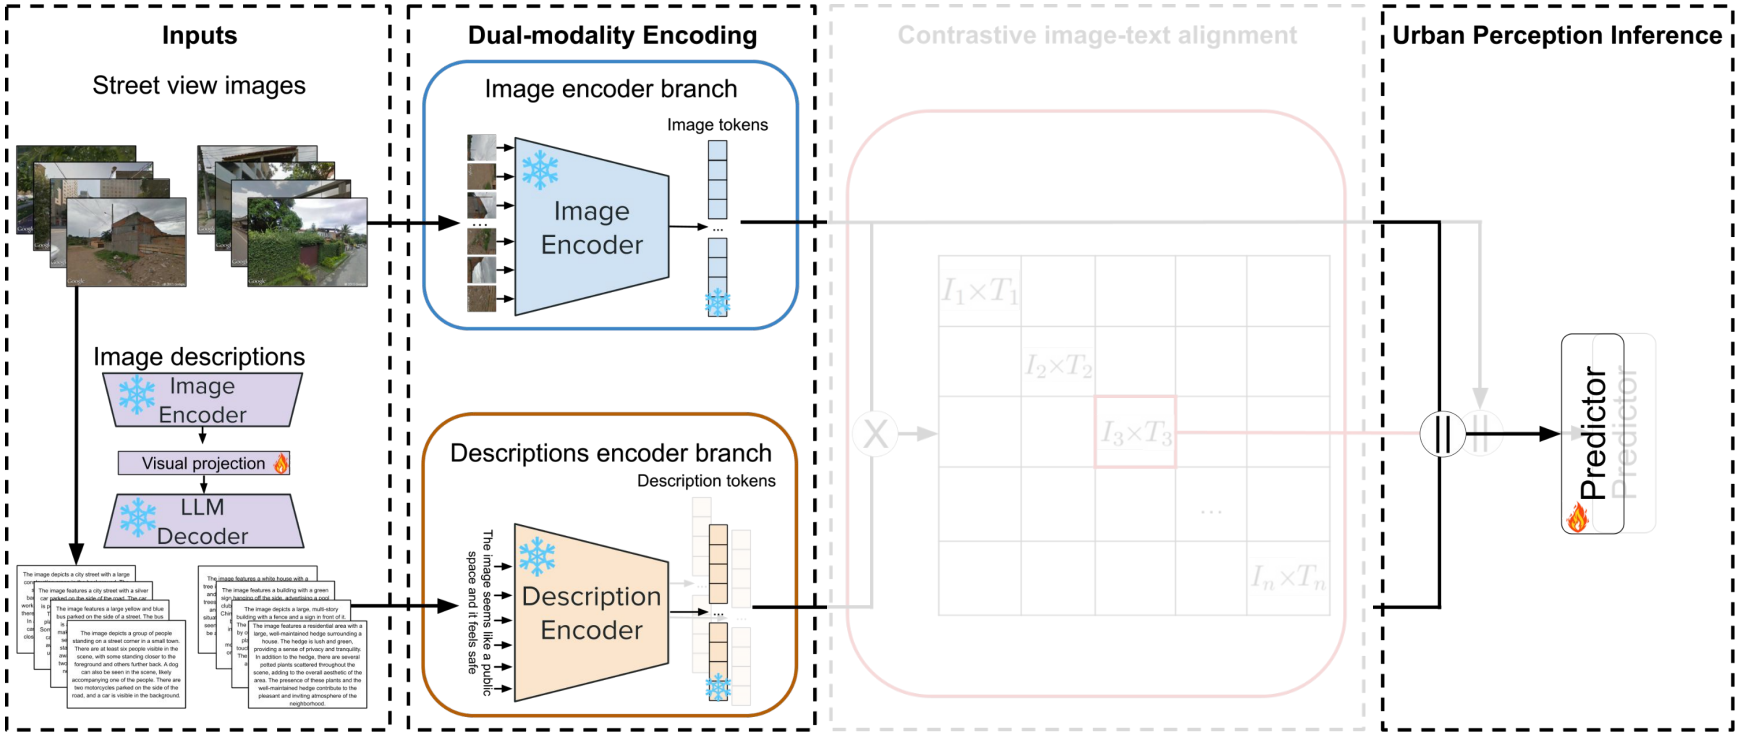

Improve the corresponding **positive** description and concatenate.

Contrastive *(learns to match image-text)*

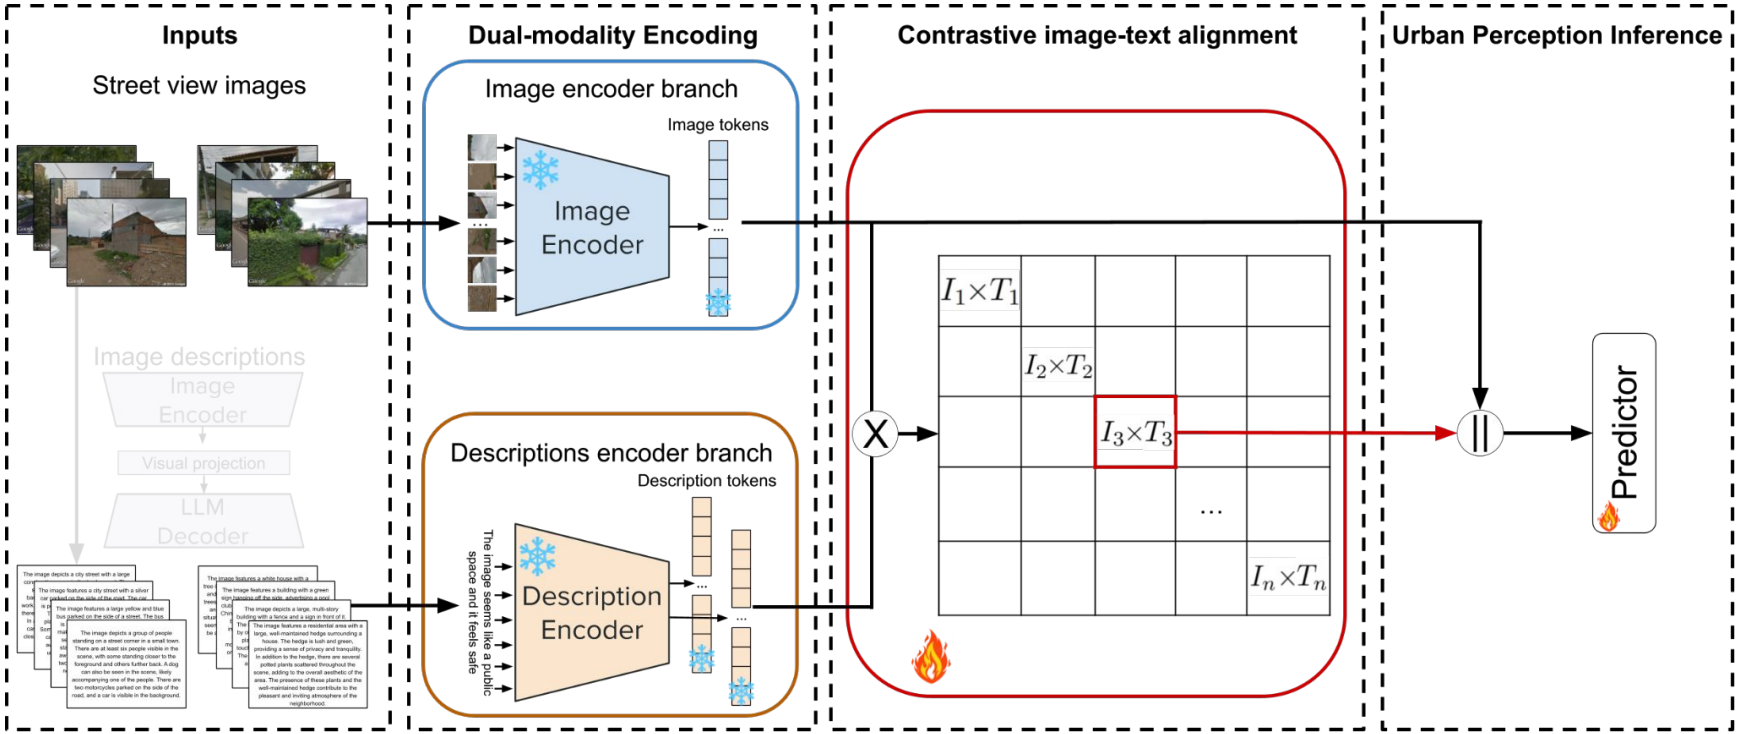

Find the best match image- **positive** and **negative** descriptions, and concatenate.

UrbanVLM *(learns all together)*

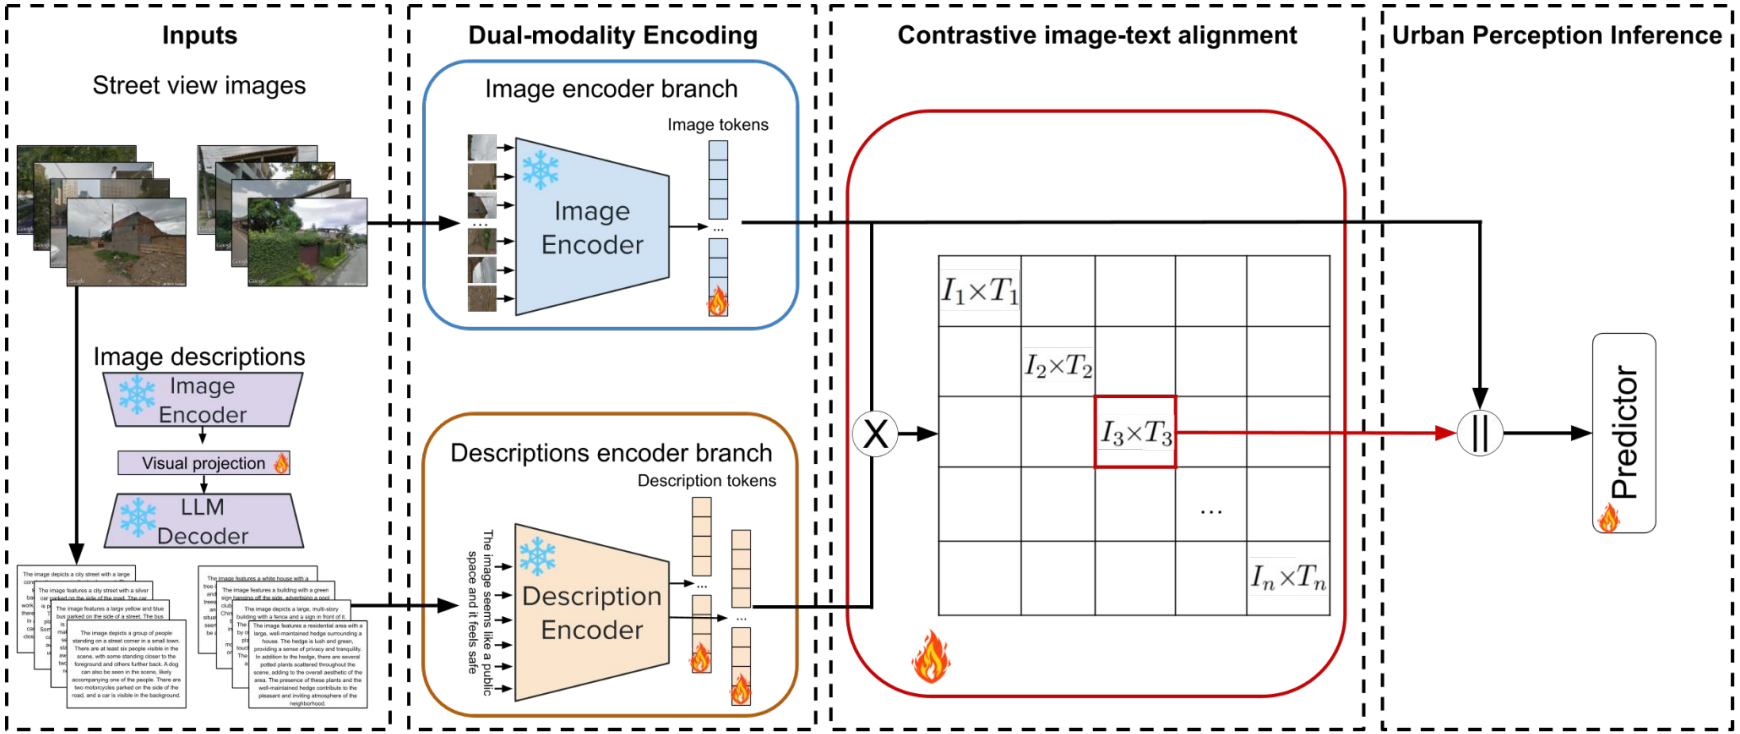

Improve and match image-text descriptions and concatenate.

Conclusions

Conclusions

- **Ablation studies** allows to analyze and understand the relevance of each component in our model.
- **Adding robust descriptions** improve the urban perception inference of images giving a human-based perception descriptions.
- **UrbanVLM** successfully learns each component and improve the classification and regression tasks by using image descriptions.

THANKS!

FGV EMAp V-DS FAPESP SÃO PAULO RESEARCH FOUNDATION

Assessing Urban Environments with Vision-Language Models

A **Comparative Analysis** of **AI-Generated Ratings** and **Human Evaluations**

Felipe Moreno-Vera and Jorge Poco

Fundação Getulio Vargas / Visual Data Science Lab
www.visualdslab.com

Ablation results

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Zero-shot	CLIP	0.39	0.41	0.39	0.24	-14.05	4.53	4.89
	SigLIP	0.57	0.43	0.57	0.45	-14.17	4.61	4.77
Only heads W/o description, contrastive & dual-modality	LlaVA+CLIP	**0.67**	**0.67**	0.66	**0.66**	**0.57**	**2.43**	**2.56**
	LlaVA+SigLIP	0.66	0.66	**0.67**	**0.66**	0.56	**2.43**	2.68
	BLIP-2+CLIP	0.63	0.61	0.62	0.61	0.53	3.4	3.21
	BLIP-2+SigLIP	0.64	0.63	0.63	0.63	0.53	3.38	3.35
Contrastive W/o description & dual-modality	LlaVA+CLIP	0.7	0.69	0.68	0.68	**0.62**	**1.81**	1.95
	LlaVA+SigLIP	**0.71**	**0.71**	**0.7**	**0.7**	0.61	1.98	**1.84**
	BLIP-2+CLIP	0.68	0.67	0.68	0.67	0.56	2.75	2.35
	BLIP-2+SigLIP	0.69	0.68	0.69	0.68	0.55	2.68	2.2
Visual projections W/o contrastive & dual-modality	LlaVA+CLIP	**0.73**	**0.72**	**0.71**	**0.71**	**0.67**	1.69	1.73
	LlaVA+SigLIP	0.72	**0.72**	**0.71**	**0.71**	0.65	**1.68**	**1.71**
	BLIP-2+CLIP	0.7	0.7	0.69	0.69	0.59	1.95	2.06
	BLIP-2+SigLIP	0.71	0.71	0.7	0.7	0.59	1.88	1.94
Dual-modality W/o description & contrastive	LlaVA+CLIP	**0.76**	**0.76**	**0.75**	**0.75**	**0.78**	1.33	**1.42**
	LlaVA+SigLIP	0.75	0.75	0.74	0.74	0.75	**1.29**	1.51
	BLIP-2+CLIP	0.72	0.72	0.73	0.72	0.69	1.6	1.34
	BLIP-2+SigLIP	0.73	0.73	0.72	0.72	0.68	1.4	1.21
UrbanVLM	LlaVA+CLIP	0.82	0.78	0.77	0.77	**0.84**	**1.04**	**0.78**
	LlaVA+SigLIP	**0.83**	**0.79**	**0.78**	**0.78**	0.83	1.08	0.79
	BLIP-2+CLIP	0.77	0.76	0.77	0.76	0.76	1.32	1.15
	BLIP-2+SigLIP	0.76	0.78	0.77	0.76	0.75	1.26	1.01

Ablation results

Ablation Study	Model Tested	Classification F-1	Regression R^2
Zero-shot	CLIP	0.24	-14.05
	SigLIP	0.45	-14.17
Only heads W/o description, contrastive & dual-modality	LlaVA+CLIP	**0.66**	**0.57**
	LlaVA+SigLIP	**0.66**	0.56
	BLIP-2+CLIP	0.61	0.53
	BLIP-2+SigLIP	0.63	0.53
Contrastive W/o description & dual-modality	LlaVA+CLIP	0.68	**0.62**
	LlaVA+SigLIP	**0.7**	0.61
	BLIP-2+CLIP	0.67	0.56
	BLIP-2+SigLIP	0.68	0.55
Visual projections W/o contrastive & dual-modality	LlaVA+CLIP	**0.71**	**0.67**
	LlaVA+SigLIP	**0.71**	0.65
	BLIP-2+CLIP	0.69	0.59
	BLIP-2+SigLIP	0.7	0.59
Dual-modality W/o description & contrastive	LlaVA+CLIP	**0.75**	**0.78**
	LlaVA+SigLIP	0.74	0.75
	BLIP-2+CLIP	0.72	0.69
	BLIP-2+SigLIP	0.72	0.68
UrbanVLM	LlaVA+CLIP	0.77	**0.84**
	LlaVA+SigLIP	**0.78**	0.83
	BLIP-2+CLIP	0.76	0.76
	BLIP-2+SigLIP	0.76	0.75

Zeroshot test

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Zero-shot	CLIP	0.39	0.41	0.39	0.24	-14.05	4.53	4.89
	SigLIP	0.57	0.43	0.57	0.45	-14.17	4.61	4.77

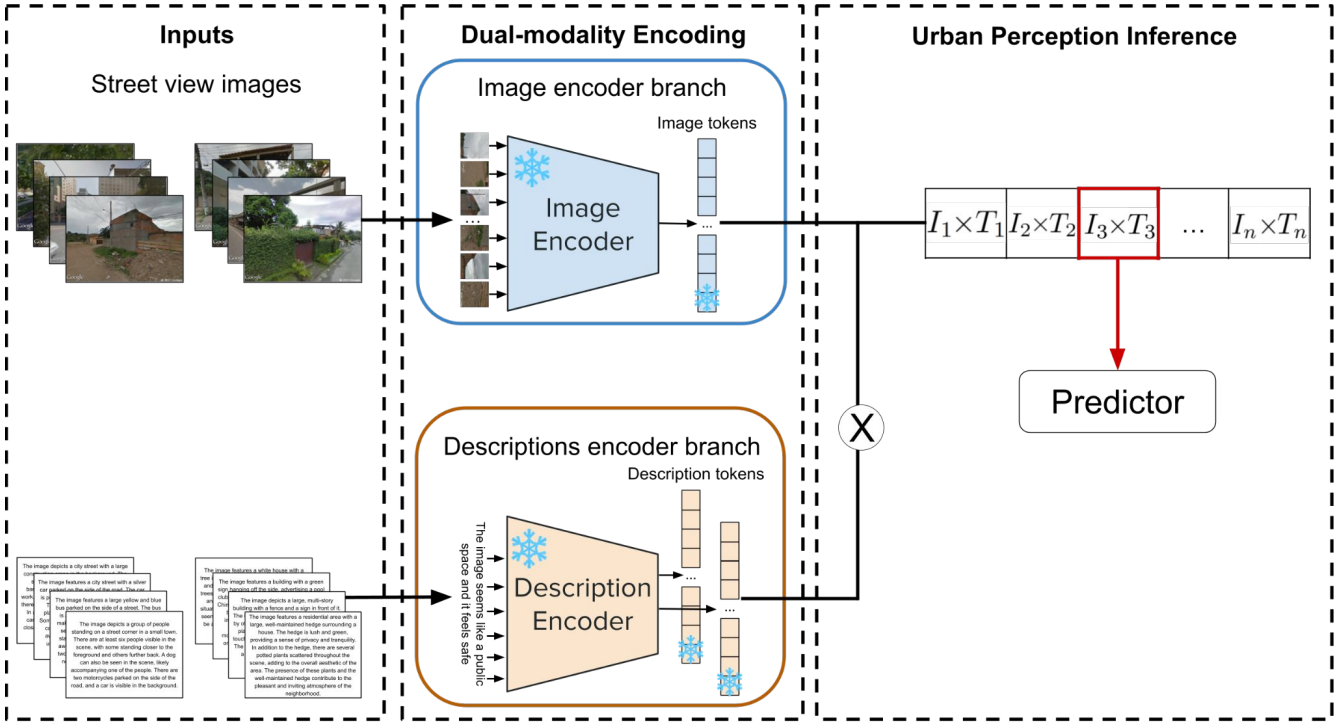

Only heads

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Only heads	LlaVA+CLIP	**0.67**	**0.67**	0.66	**0.66**	**0.57**	**2.43**	**2.56**
W/o description, contrastive & dual-modality	LlaVA+SigLIP	0.66	0.66	**0.67**	**0.66**	0.56	**2.43**	2.68
	BLIP-2+CLIP	0.63	0.61	0.62	0.61	0.53	3.4	3.21
	BLIP-2+SigLIP	0.64	0.63	0.63	0.63	0.53	3.38	3.35

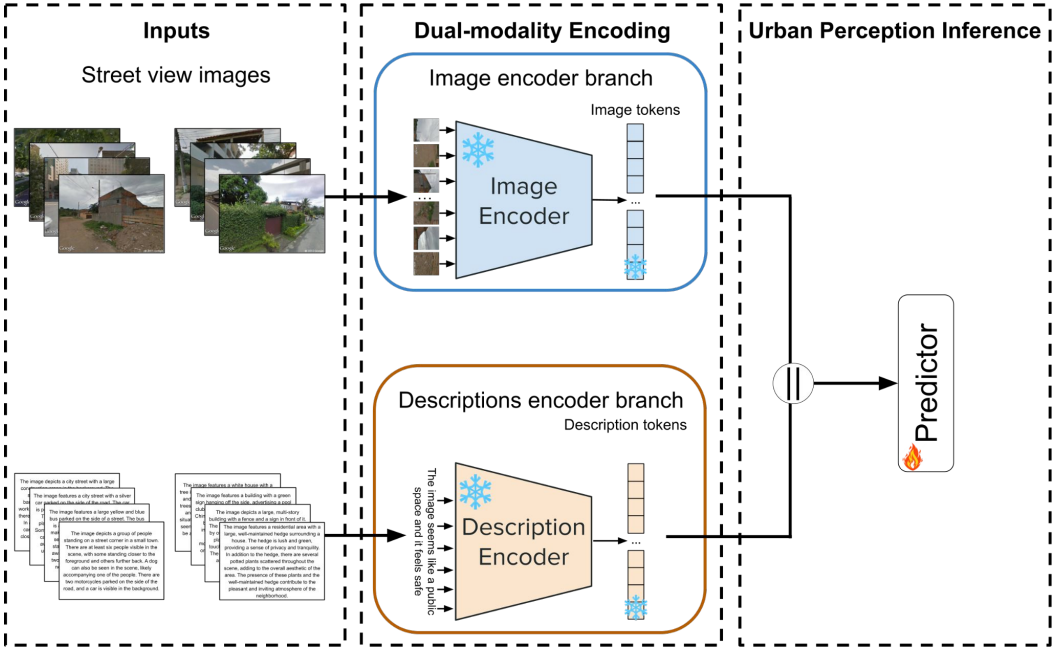

Contrastive

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Contrastive	LlaVA+CLIP	0.7	0.69	0.68	0.68	**0.62**	**1.81**	1.95
W/o description & dual-modality	LlaVA+SigLIP	**0.71**	**0.71**	**0.7**	**0.7**	0.61	1.98	**1.84**
	BLIP-2+CLIP	0.68	0.67	0.68	0.67	0.56	2.75	2.35
	BLIP-2+SigLIP	0.69	0.68	0.69	0.68	0.55	2.68	2.2

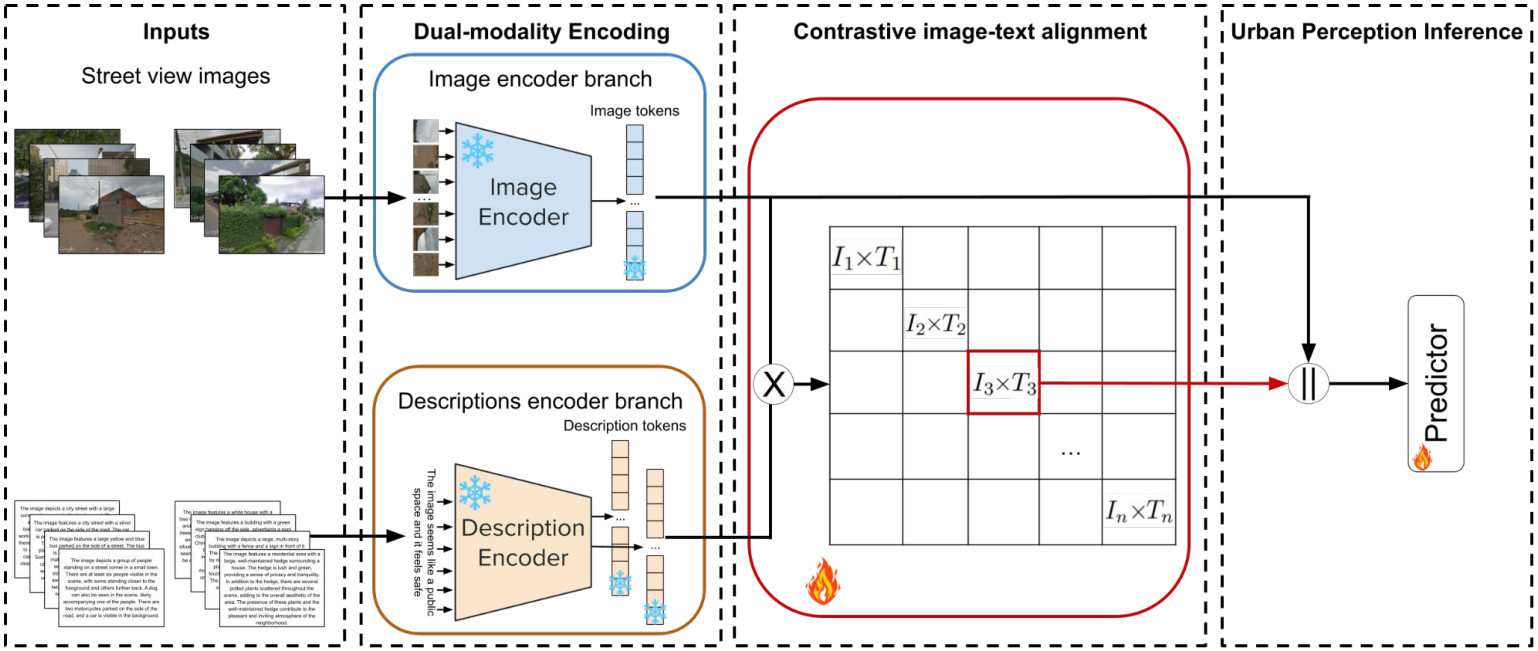

Visual projections

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Visual projections	LlaVA+CLIP	**0.73**	**0.72**	**0.71**	**0.71**	**0.67**	1.69	1.73
W/o contrastive & dual-modality	LlaVA+SigLIP	0.72	**0.72**	**0.71**	**0.71**	0.65	**1.68**	**1.71**
	BLIP-2+CLIP	0.7	0.7	0.69	0.69	0.59	1.95	2.06
	BLIP-2+SigLIP	0.71	0.71	0.7	0.7	0.59	1.88	1.94

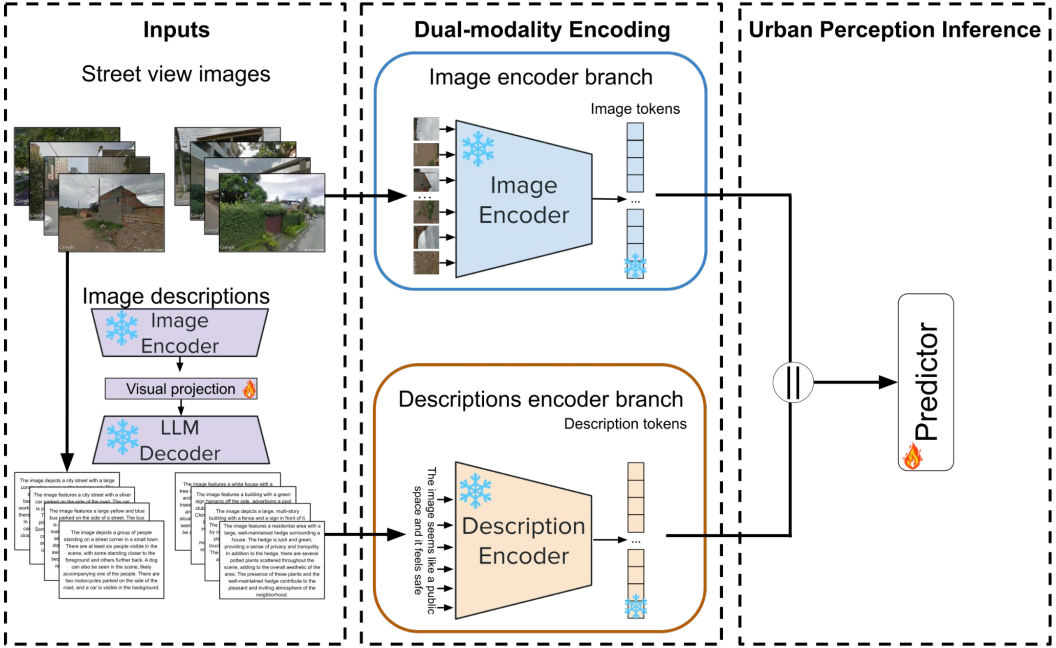

Dual-modality

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
Dual-modality	LlaVA+CLIP	**0.76**	**0.76**	**0.75**	**0.75**	**0.78**	1.33	**1.42**
W/o description & contrastive	LlaVA+SigLIP	0.75	0.75	0.74	0.74	0.75	**1.29**	1.51
	BLIP-2+CLIP	0.72	0.72	0.73	0.72	0.69	1.6	1.34
	BLIP-2+SigLIP	0.73	0.73	0.72	0.72	0.68	1.4	1.21

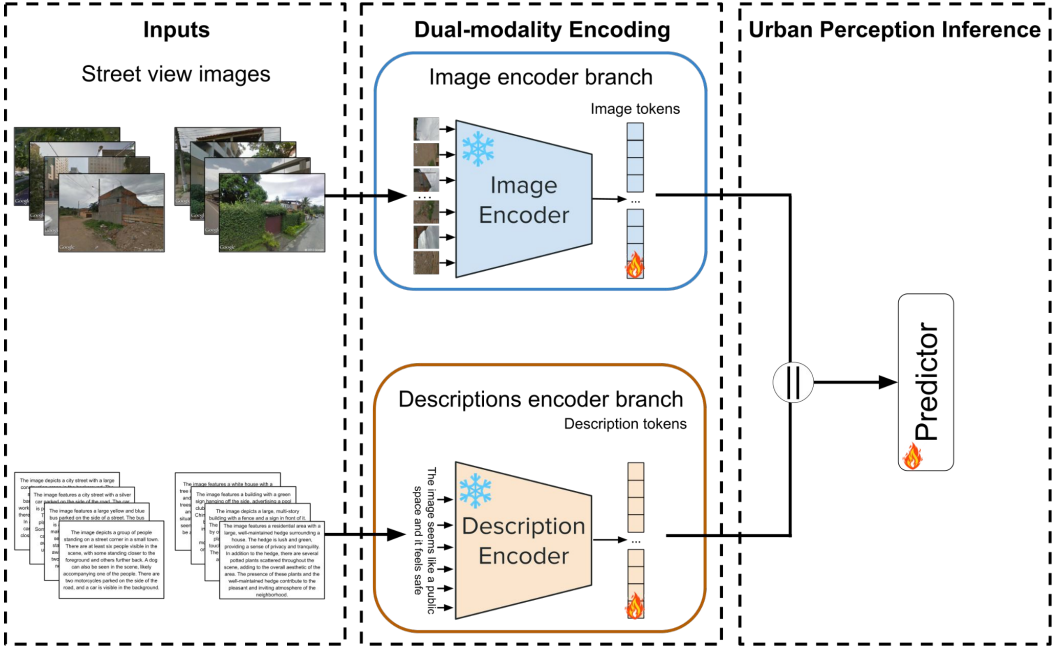

UrbanVLM

Ablation Study	Model Tested	Classification				Regression		
		Acc	Precision	Recall	F-1	R^2	RMSE	MAE
UrbanVLM	LlaVA+CLIP	0.82	0.78	0.77	0.77	**0.84**	**1.04**	**0.78**
	LlaVA+SigLIP	**0.83**	**0.79**	**0.78**	**0.78**	0.83	1.08	0.79
	BLIP-2+CLIP	0.77	0.76	0.77	0.76	0.76	1.32	1.15
	BLIP-2+SigLIP	0.76	0.78	0.77	0.76	0.75	1.26	1.01

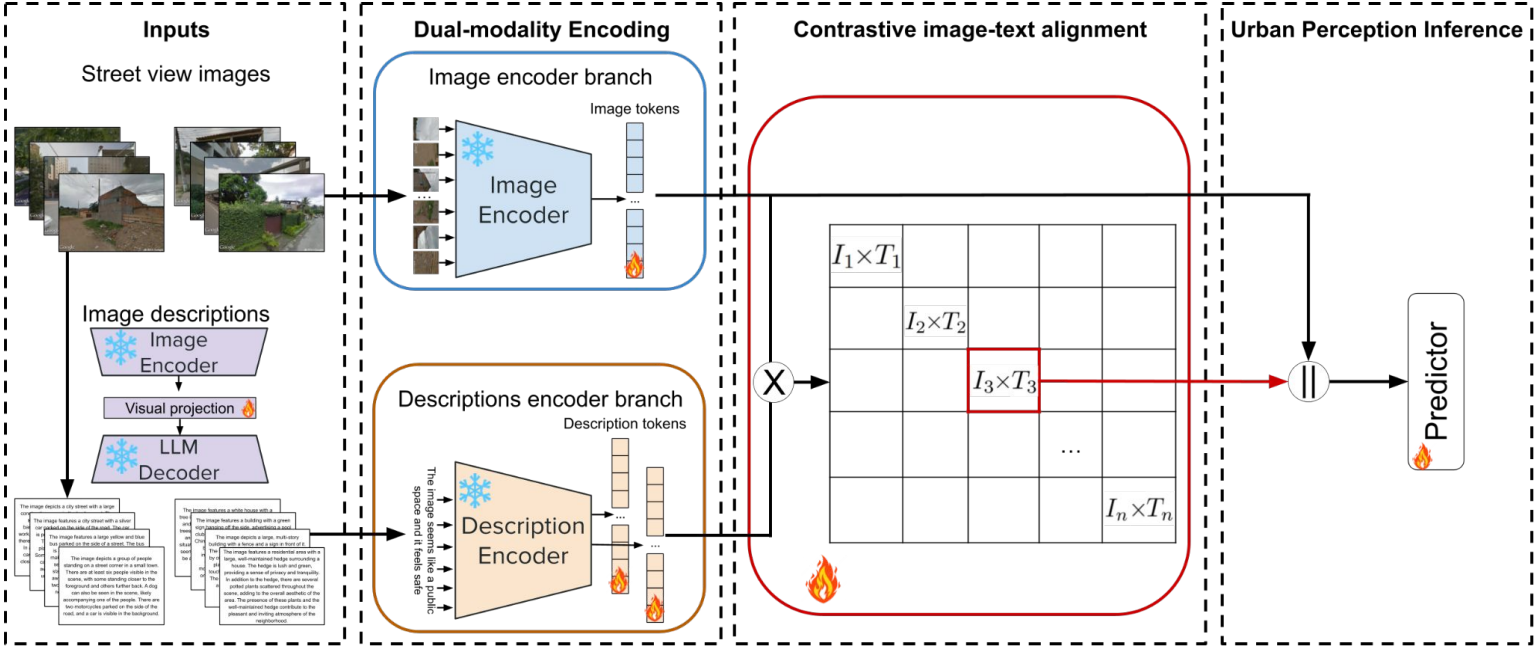

Scores correlation

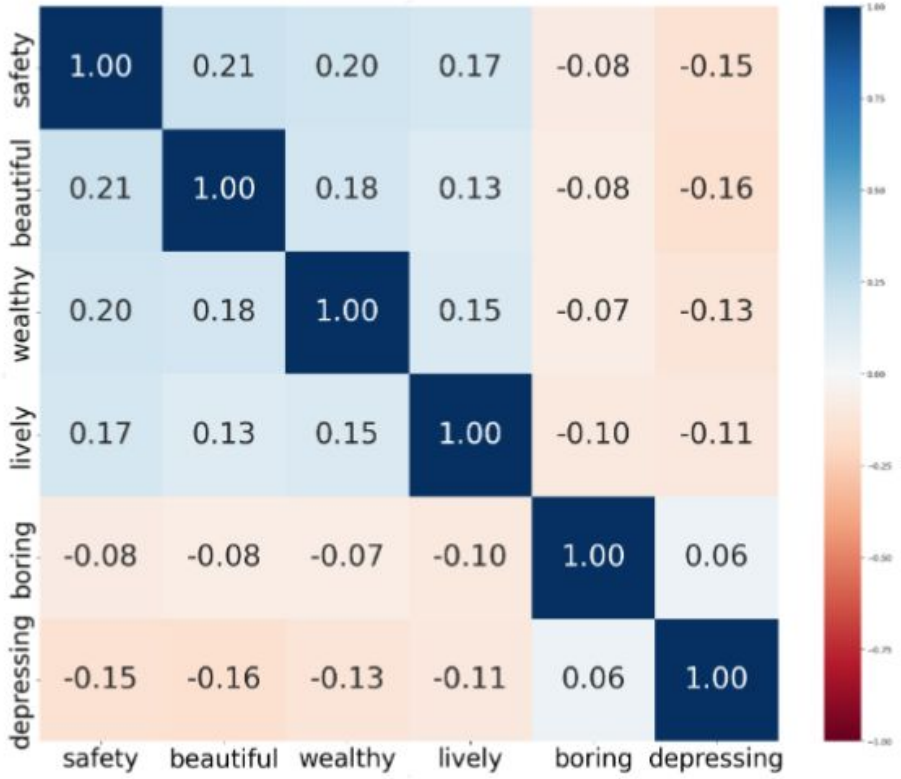

(a) "Strength of schedule" algorithm

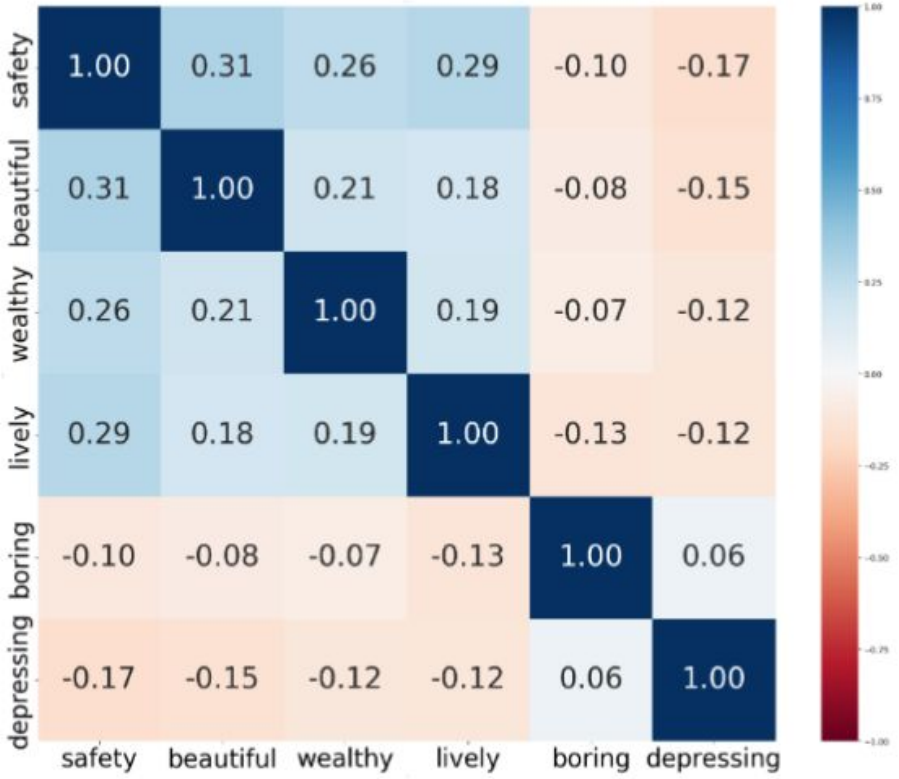

(b) UrbanVLM predictions